

Project 1: D.C. crime

TOTAL: 10pts

PURPOSE

To analyze crime in Washington, D.C. This project introduces you to connecting to data, creating new fields, and conducting trend analysis.

GOAL

Identify crime trends in Washington, DC over behavior, space, and time.

DATA

This project uses real live Washington, D.C. crime data. [It's located here.](#)

METHODS

Step 0: Load and install packages

1. You will need two additional packages to install and load into your script:
 - a. `install.packages('tidyverse')` # run this once, ever
 - b. `library(tidyverse)` # run this every session
 - c. `install.packages('explore')` # run this once, ever
 - d. `library(explore)` # run this every session
 - e. `install.packages('GGally')` # run this once, ever
 - f. `library(GGally)` # run this every session

Step 1: Examine data

1. Take a look at the table - get comfortable with the fields captured, and how it is organized.
2. [Look at the 2020 D.C. crime data here.](#)
3. Take a look at the table - examine the fields captured, and how it is organized.

Step 2: Read data

1. In the top left pane in RStudio, add this line of code:
 - a. `dc.data2020 <- read.csv("", stringsAsFactors = FALSE)`
 - b. In between the quotes, look for the table ID found as part of the URL for the data source.
 - c. It should look something like this:
 - i. `dc.data2020 <-
read.csv("https://opendata.arcgis.com/datasets/f516e0dd7b614b088ad781b0c4002331_2.csv", stringsAsFactors = FALSE)`
 - d. Highlight that line of code and Run it.
 - e. Success!
 - f. In the top right pane, click on the output to view it. Should be a table of all 2020 D.C. crimes to date. Close that table tab after viewing.
2. Repeat this process for 2021.
 - a. From the data website, edit the URL.
 - i. Huh?
 - ii. This is the URL: <https://opendata.dc.gov/datasets/crime-incidents-in-2020/data>
 - iii. Change the '2020' to '2021'
 - b. Click on 'I want to use this' at the bottom left
 - c. Click on 'View API Resources'
 - d. On GeoJSON, click the 'copy link' button
 - e. In the top left pane in RStudio, add this line of code:
 - i. `dc.data2021 <- read.csv("", stringsAsFactors = FALSE)`
 - ii. In between the quotes, paste the link you copied from the data website, and change the .json extension at the end to .csv
 - iii. Highlight that line of code and Run it.
 - iv. More success! A second table added to your environment.
3. Repeat the process from #2 for every year since 2008.
 - a. When complete, you'll have yearly tables in your Environments pane.
 - b. Like this:
 - i. `dc.data2025 <-
read.csv("https://opendata.arcgis.com/datasets/74d924ddc3374e3b977e6f002478cb9b_7.csv", stringsAsFactors = FALSE)`
 - ii. `dc.data2024 <-
read.csv("https://opendata.arcgis.com/datasets/c5a9f33ffca546babbd91de1969e742d_6.csv", stringsAsFactors = FALSE)`
 - iii. `dc.data2023 <-
read.csv("https://opendata.arcgis.com/datasets/89561a4f02ba46cca3c42333425d1b87_5.csv", stringsAsFactors = FALSE)`
 - iv. `dc.data2022 <-
read.csv("https://opendata.arcgis.com/datasets/f9cc541fc8c04106a05a1a4f1e7e813c_4.csv", stringsAsFactors = FALSE)`
 - v. `dc.data2021 <-
read.csv("https://opendata.arcgis.com/datasets/619c5bd17ca2411db0689bb0a211783c_3.csv", stringsAsFactors = FALSE)`

```

vi.    dc.data2020 <-
      read.csv("https://opendata.arcgis.com/datasets/f516e0dd7b614b088ad781b0c40
      02331_2.csv", stringsAsFactors = FALSE)
vii.   dc.data2019 <-
      read.csv("https://opendata.arcgis.com/datasets/f08294e5286141c293e9202fcd3e
      8b57_1.csv", stringsAsFactors = FALSE)
viii.  dc.data2018 <-
      read.csv("https://opendata.arcgis.com/datasets/38ba41dd74354563bce28a359b
      59324e_0.csv", stringsAsFactors = FALSE)
ix.    dc.data2017 <-
      read.csv("https://opendata.arcgis.com/datasets/6af5cb8dc38e4bcbac8168b27ee
      104aa_38.csv", stringsAsFactors = FALSE)
x.     dc.data2016 <-
      read.csv("https://opendata.arcgis.com/datasets/bda20763840448b58f8383bae80
      0a843_26.csv", stringsAsFactors = FALSE)
xi.    dc.data2015 <-
      read.csv("https://opendata.arcgis.com/datasets/35034fcb3b36499c84c94c069ab
      1a966_27.csv", stringsAsFactors = FALSE)
xii.   dc.data2014 <-
      read.csv("https://opendata.arcgis.com/datasets/6eaf3e9713de44d3aa103622d51
      053b5_9.csv", stringsAsFactors = FALSE)
xiii.  dc.data2013 <-
      read.csv("https://opendata.arcgis.com/datasets/5fa2e43557f7484d89aac9e1e76
      158c9_10.csv", stringsAsFactors = FALSE)
xiv.   dc.data2012 <-
      read.csv("https://opendata.arcgis.com/datasets/010ac88c55b1409bb67c9270c8f
      c18b5_11.csv", stringsAsFactors = FALSE)
xv.    dc.data2011 <-
      read.csv("https://opendata.arcgis.com/datasets/9d5485ffae914c5f97047a7dd86e
      115b_35.csv", stringsAsFactors = FALSE)
xvi.   dc.data2010 <-
      read.csv("https://opendata.arcgis.com/datasets/fdacfbdda7654e06a161352247d
      3a2f0_34.csv", stringsAsFactors = FALSE)
xvii.  dc.data2009 <-
      read.csv("https://opendata.arcgis.com/datasets/73cd2f2858714cd1a7e2859f8e6e
      4de4_33.csv", stringsAsFactors = FALSE)
xviii. dc.data2008 <-
      read.csv("https://opendata.arcgis.com/datasets/180d56a1551c4e76ac2175e63dc
      0dce9_32.csv", stringsAsFactors = FALSE)

```

Step 3: Merge data

Combine all of the yearly tables into one with this:

- `data.temp <- rbind(dc.data2008, dc.data2009, dc.data2010, dc.data2011, dc.data2012, dc.data2013, dc.data2014, dc.data2015, dc.data2016, dc.data2017, dc.data2018, dc.data2019, dc.data2020, dc.data2021, dc.data2022, dc.data2023, dc.data2024, dc.data2025)`

Step 4: Clean data

Parse the 'REPORT_DAT' field into separate data and time fields:

- `dc.data <- separate(data.temp, REPORT_DAT, into = c("date", "time"), sep = " ")`

Step 5: Create data

Create six usable fields (columns) in your data:

1. Format the date column
 - a. `dc.data$date <- as.Date(dc.data$date, format = "%Y/%m/%d")`
2. Create an hour of day column
 - a. `dc.data$hour <- substr(dc.data$time, 0, 2)`
 - b. `dc.data$hour <- as.numeric(dc.data$hour)`
3. Create day of week
 - a. `dc.data$dow <- weekdays(dc.data$date)`
4. Create week of year
 - a. `dc.data$week <- format(as.Date(dc.data$date), "%U")`
 - b. `dc.data$week <- as.numeric(dc.data$week)`
5. Create month of year
 - a. `dc.data$months <- months(dc.data$date)`
6. Create year
 - a. `dc.data$year <- substr(dc.data$date, 0, 4)`
7. Seasons?

Step 6: Analyze data

1. Explore your data.
 - a. `explore(dc.data)`
2. You'll get a new pop-up window.
 - a. Change the 'variable' to different fields to analyze them.
 - b. Also use the 'target' menu to add complexity.
 - c. Try different variables.
 - d. Try combinations of variables.
 - e. Take some time and get loose. Identify trends, which are common themes across the data.
 - f. Close this window when complete (your RStudio session stays 'active' until you close it)

[HERE IS THE SCRIPT WE BUILT IN CLASS ON 22 JANUARY 2025](#)

Step 7: Visualize data

1. Here's a (relatively) quick visualization for the numeric fields in these data:
 - a. `data.new <- select(dc.data, hour, DISTRICT, WARD, week, year)`
 - b. `ggpairs(data.new) + theme_bw()`
 - i. But this gives you an error, because there are more than 15 unique year values... so we have to adjust the 'cardinality' parameter, like this:
 1. `ggpairs(data.new, cardinality_threshold = xx) + theme_bw()`

Step 8: Intro to graphing data

Bar graphs:

- Follow these iterative steps to make progressively better bar graphs:
 - `ggplot()`
 - `ggplot(dc.data, aes(SHIFT))`
 - `ggplot(dc.data, aes(SHIFT)) + geom_bar(stat = "count")`
 - `ggplot(dc.data, aes(SHIFT)) + geom_bar(stat = "count", aes(fill = OFFENSE))`
 - `ggplot(dc.data, aes(SHIFT)) + geom_bar(stat = "count", aes(fill = OFFENSE), position = position_stack(reverse = TRUE))`
 - `ggplot(dc.data, aes(SHIFT)) + geom_bar(stat = "count", aes(fill = OFFENSE), position = position_stack(reverse = TRUE)) + theme_classic()`
 - `ggplot(dc.data, aes(SHIFT)) + geom_bar(stat = "count", aes(fill = OFFENSE), position = position_stack(reverse = TRUE)) + theme_classic() + theme(legend.position = "bottom")`
 - `ggplot(dc.data, aes(SHIFT)) + geom_bar(stat = "count", aes(fill = OFFENSE), position = position_stack(reverse = TRUE)) + theme_classic() + facet_wrap(~DISTRICT)`

Formatting:

- Make your graph an object, and then you can customize the formatting:
 - `ggplot(dc.data, aes(SHIFT)) + geom_bar(stat = "count", aes(fill = OFFENSE), position = position_stack(reverse = TRUE)) + theme_classic() +`
 - `labs(`
 - `title = "Your title here",`
 - `x = "x axis label here",`
 - `y = "y axis label here"`
 - `) +`
 - `theme(`
 - `panel.grid.minor = element_blank(),`
 - `panel.grid.major.y = element_blank(),`
 - `panel.grid.major.x = element_line(),`
 - `axis.ticks = element_blank()`
 - `)`

Better graphs:

These aren't mandatory, but a little cleaner and better than the examples above.

- First, some data prep to group crime by 'Person' and 'Property':
 - `dc.data$TYPE <- case_when(`
 - `dc.data$OFFENSE %in%`
 - `c(`
 - `"ARSON",`
 - `"BURGLARY",`
 - `"MOTOR VEHICLE THEFT",`
 - `"THEFT F/AUTO",`
 - `"THEFT/OTHER") ~ "Property",`
 - `dc.data$OFFENSE %in%`
 - `c("ASSAULT W/DANGEROUS WEAPON", "HOMICIDE", "ROBBERY", "SEX ABUSE")`
 - `~ "Person")`
- Second, a cleaner crimes by hour graph:
 - `dc.data %>% ggplot() +`
 - `geom_line(aes(x = xxxx), stat = "count", group = 1, color = "blue", size = 1) +`
 - `labs(title = "DC Crimes by Hour Reported \n2008 - 202x", x = "Hour Reported", y =`
 - `"Number of Crimes (Thousands)", fill = "Type of Crime") +`
 - `theme(plot.title = element_text(hjust = 0.5, size = 14),`
 - `axis.text.x = element_text(color = "black", size = 10),`
 - `axis.text.y = element_text(color = "black", size = 10),`
 - `axis.ticks.y = element_blank(),`
 - `panel.grid.major.x = element_blank(),`
 - `panel.grid.major.y = element_line(color = "gray"),`
 - `panel.background = element_blank()) +`
 - `scale_y_continuous(limits = c(0, xxxxx), breaks = c(0,`
 - `5000,10000,15000,20000,25000,30000,35000), labels = c(0,5,10,15,20,25,30,35))`
- Third, graph crime type by district:
 - `dc.data %>% filter(DISTRICT != "NA") %>%`
 - `ggplot() +`
 - `geom_bar(aes(x = as.factor(xxx), fill = TYPE), stat = "count") +`
 - `labs(title = "Crimes by DC Police District \n2010 - 202x", x = "District", y = "Number of`
 - `Crimes (Thousands)", fill = "Type of Crime") +`
 - `theme(plot.title = element_text(hjust = 0.5, size = 14),`
 - `axis.text.x = element_text(color = "black", size = 10),`
 - `axis.text.y = element_text(color = "black", size = 10),`
 - `axis.ticks.y = element_blank(),`
 - `panel.grid.major.x = element_blank(),`
 - `panel.grid.major.y = element_line(color = "gray"),`
 - `panel.background = element_blank()) +`
 - `scale_y_continuous(limits = c(0, 125000), breaks = c(0,`
 - `25000,50000,75000,100000,125000), labels = c(0,25,50,75,100,125))`

- Fourth, graph a crime type as a percentage of weapon type/method:
 - `dc.data %>% filter(TYPE == "xxx") %>%`
 - `ggplot() +`
 - `geom_bar(aes(x = OFFENSE, fill = METHOD), position = "fill") +`
 - `labs(title = "DC xxx Crimes by Method \n2010 - 202x", x = "Crime", y = "Proportion of Crime", fill = "Method") +`
 - `theme(plot.title = element_text(hjust = 0.5, size = 14),`
 - `axis.text.x = element_text(color = "black", size = 10),`
 - `axis.text.y = element_text(color = "black", size = 10),`
 - `axis.ticks.y = element_blank(),`
 - `panel.grid.major.x = element_blank(),`
 - `panel.grid.major.y = element_line(color = "gray"),`
 - `panel.background = element_blank()) +`
 - `scale_x_discrete(labels = c("Assault", "Homicide", "Robbery", "Sex Abuse")) +`
 - `scale_y_continuous(labels = c("0%", "25%", "50%", "75%", "100%")) +`
 - `scale_fill_manual(values = c("#1b9e77", "#d95f02", "#7570b3"), labels = c("Gun", "Knife", "Other"))`

[HERE IS THE SCRIPT WE CREATED IN CLASS ON 29 JANUARY 2025](#)

ANALYSIS

When conducting your analysis, consider this perspective:

1. State a finding
 - a. This is the new thing that you've identified from the data
2. Provide your evidence
 - a. This is typically a statistical reference
3. Add the context
 - a. This is the integration of the stat(s) and the finding
 - b. Compare your finding to the norm/average, or a similar measure in the data

FORMAT

Submit your script as a usable HTML file, via RMarkdown. You create this output as a .Rmd file in RStudio, but only submit the .html file.

[Use this file as a template.](#)

[HERE IS THE TEMPLATE WE CREATED IN CLASS ON 29 JANUARY 2025](#)

[HERE IS THE TEMPLATE WE FURTHER REFINED IN CLASS ON 05 FEBRUARY 2025](#)

SUBMISSION

Once your analysis is complete, please submit your project as either an HTML file, via Canvas.

GRADES

Analysis (4pts)

- Identify four unique analytic findings (*1pt each*)
 - Each finding should be 3-4 sentences, and describe something *specific* about the data.
 - *Each find should include a specific statistical reference*

Visuals (4pts)

- Include two visuals (*2pts each*)
 - These visuals should be thoughtful and logical (make sense to someone that has never seen them before); somehow, somehow referenced in the text; and analytically meaningful (provide useful, relevant, and actionable insights)

Grammar (1pt)

- Error-free writing. No typos, run-on sentences, or sentence fragments. Proper punctuation. Real words.

File format (1pt)

- Project, including all analysis and visualizations, created in RMarkdown and submitted as an HTML file via Canvas. File is neatly and cleanly formatted, hiding unnecessary code as appropriate.

Please [email me](#) with any questions.